Communism in Western Political Discourse

Yu Liang
Department of Computer Science, Duke University
April 2026

Abstract

This research studies how the meaning of *communism* is constructed across Western political and intellectual texts from 1931 to 1990. Instead of treating *communism* as a fixed keyword, I analyze the complete sentences in which it appears. I combine distributional analysis, spaCy named entity recognition, BERT contextual embeddings, weak supervision, logistic regression, and close reading. The analysis shows that *communism* is structured by two interacting dimensions. First, unsupervised BERT clustering separates the corpus into two contextual frames: DOCTRINAL IMAGINATION, where *communism* appears as a belief system, moral philosophy, historical theory, or quasi-religious worldview; and GEOPOLITICAL MOBILIZATION, where it appears as a political movement, state project, international force, or Cold War actor. Second, weakly supervised classification shows that these frames are crossed by stance: sentences position *communism* as HOSTILE, NEUTRAL, or SYMPATHETIC. The central finding is that Western discourse does not assign *communism* one stable meaning. It makes the concept legible by shifting between what kind of object communism is imagined to be and what evaluative stance the text takes toward it.

1 Introduction

Political concepts rarely behave like ordinary keywords. A politically loaded word such as *communism* does not simply refer to a doctrine; it can name a belief system, a future society, a foreign danger, a party organization, a revolutionary project, or a moral challenge. This makes *communism* an ideal case for asking a broader text-as-data question: **how is the meaning of communism constructed and varied across texts?**

The problem is that different computational methods see different parts of meaning. A frequency count can show where the word appears, but not what it does. Named entity recognition can identify people, places, and organizations around the word, but not fully explain how the concept is being framed. BERT embeddings can reveal contextual structure, but their clusters require interpretation. Supervised classification can test whether a proposed interpretation is predictable, but it depends on the quality of labels.

This paper turns that methodological tension into the structure of the analysis. I first extract every complete sentence containing *communism*. I then move from surface distribution to named entities, from named entities to BERT contextual frames, and from BERT frames to weakly supervised stance classification. The resulting story is layered: *communism* is not merely anti-communist, pro-communist, or neutral. It is constructed through a deeper contrast between DOCTRINAL IMAGINATION and GEOPOLITICAL MOBILIZATION, and then evaluated through hostile, neutral, or sympathetic stance.

2 Related Work

This project draws on several strands of NLP and computational social science. First, the basic distributional idea that meaning is shaped by usage follows a long linguistic tradition often summarized by the claim that words are known by the company they keep [Firth, 1957]. Modern distributional semantics operationalizes

Field	Value
Target concept	*communism*
Unit of analysis	Sentence containing *communism*
Total concept sentences	227
Unique documents	23
Unique sources	4
Corpus buckets	3
Year range	1931–1990

Table 1: Dataset summary after extracting complete sentences containing *communism*.

this intuition through vector representations such as word2vec and GloVe [Mikolov et al., 2013, Pennington et al., 2014]. These methods motivate the central assumption of this project: the meaning of *communism* should be studied through its contexts, not through dictionary definition alone.

Second, text-as-data work emphasizes both the promise and limits of automatic content analysis for political language. Computational models can reveal patterns across corpora, but categories still require interpretation, especially when concepts are contested [Grimmer and Stewart, 2013, Roberts et al., 2014]. This is central here because *communism* is not a stable object but a historically loaded political concept.

Third, named entity recognition provides a way to identify the people, places, institutions, and groups through which concepts become anchored. Historical and political texts, however, often challenge standard NER systems because they contain shifting institutional names, ideological labels, and domain-specific expressions [Ehrmann et al., 2023, Lample et al., 2016]. I therefore use NER as a tool for interpretive support, not as a complete semantic model.

Fourth, contextual embeddings make it possible to represent the same word differently across contexts. BERT is especially useful because it encodes sentence-level context rather than assigning one vector to one word type [Devlin et al., 2019]. More generally, embedding-based work on semantic change shows that meaning can be studied through usage patterns across corpora and time [Hamilton et al., 2016, Kozlowski et al., 2019].

Finally, weak supervision provides a practical way to create training labels when full manual annotation is costly or when categories are exploratory [Ratner et al., 2020]. In this project, weak labels operationalize stance toward *communism*: hostile, neutral, or sympathetic.

3 Corpus and Research Design

The corpus contains Western political and intellectual texts that discuss communism in three settings: official anti-communist discourse, mainstream Western commentary, and Western left-internal critique. This is a useful corpus because the same concept appears in rhetorically different environments. Official speeches and platforms often tie *communism* to danger, freedom, peace, and national security. Mainstream essays often treat it as a moral, religious, or philosophical problem. Left-internal texts often discuss communist parties, revolution, capitalism, democracy, and ideological dispute.

The corpus spans 1931 to 1990, covering prewar anti-communism, early Cold War containment, postwar American political rhetoric, Eurocommunism, and the decline of Western communist parties. These historical settings make the corpus appropriate for studying variation in meaning rather than merely word frequency.

3.1 Data Preparation

The target concept is the exact word *communism*. The unit of analysis is the complete sentence containing that word. This produced 227 sentence-level contexts across 23 documents, 4 sources, and 3 corpus buckets.

Doc ID	Year	Title	Count
official_01	1956	The Meaning of Communism to Americans	60
mainstream_01	1931	The Religion of Communism	42
mainstream_02	1939	How to Tell a Communist	26
left_05	1978	Eurocommunism is Anti-Communism	14
left_01	1934	If America Should Go Communist	13
mainstream_09	1961	The Failure of Communism	10
official_07	1957	Eisenhower Doctrine	10
mainstream_05	1951	The Moral Challenge of Communism	10

Table 2: Documents with the most sentences containing *communism*. The concentration of occurrences shows that the concept is often an object of sustained argument rather than a passing reference.

The first pattern is concentration. The word is not evenly distributed. Nixon's 1956 study paper contributes 60 sentences, and "The Religion of Communism" contributes 42. These are not incidental mentions: they are texts built around defining, explaining, or contesting the concept.

The extracted sentences are distributed across three corpus buckets rather than coming from a single rhetorical environment. As Table 3 shows, mainstream Western commentary and official anti-communist discourse contribute almost the same number of concept sentences, while Western left-internal critique contributes fewer but analytically important cases. This balance matters for the project: it allows *communism* to be compared across public commentary, state-centered anti-communism, and insider left critique.

Corpus Bucket	Sentences
Mainstream Western commentary	94
Official anti-communist discourse	91
Western left-internal critique	42

Table 3: Distribution of sentences containing *communism* across corpus buckets. The corpus is not dominated by a single discourse type: mainstream commentary and official anti-communist texts are nearly balanced, while left-internal critique provides a smaller but distinctive contrast set.

4 Named Entities: Where the Concept is Anchored

The first interpretive question is whether *communism* is associated with particular people, institutions, places, or political groups. I ran spaCy NER over all concept sentences and focused on PERSON, GPE, ORG, and NORP. The result immediately shows that *communism* is politically anchored. NORP, ORG, and GPE dominate the entity distribution.

Entity Type	NORP	ORG	GPE	DATE	PERSON	CARDINAL	LOC	ORDINAL	MONEY	LAW
Count	78	52	51	22	17	16	14	9	5	2

Table 4: Named entity type frequency. The high counts for NORP, ORG, and GPE show that *communism* is not discussed as an abstract term alone; it is attached to groups, institutions, and places.

The top focus entities sharpen this point. The most common entities include *Communist*, *American*, *International Communism*, *America*, *the United States*, *Russia*, *Soviet*, and *the Soviet Union*. These are not neutral companions. They reveal the main arenas in which the concept is made meaningful: American self-definition, Soviet power, international conflict, and communist party identity.

NER is helpful because it tells us *where* the concept is anchored. But it cannot tell us *how* the concept is framed. For that, I turn to contextual embeddings.

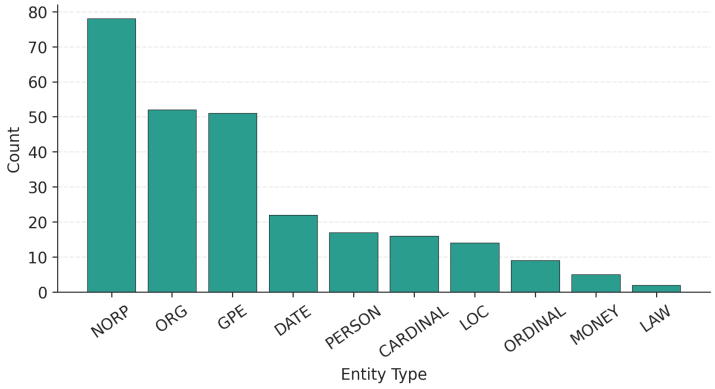

Figure 1: Named entity type frequency in sentences containing *communism*. NORP, ORG, and GPE dominate, suggesting that the concept is repeatedly constructed through political group identity, institutions, and geopolitical space.

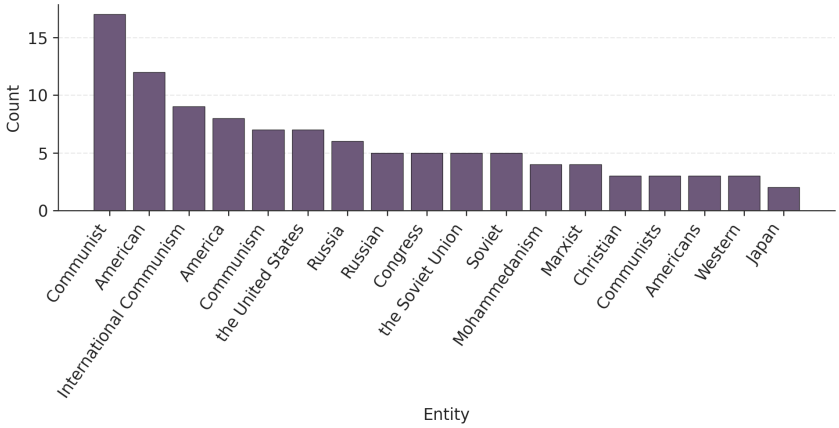

Figure 2: Top PERSON/GPE/ORG/NORP entities near *communism*. The most frequent entities point toward a discourse organized around American identity, Soviet power, international communism, and communist political identity.

Entity Type	**Entity**	**Representative Co-occurrence**
GPE	United States / Soviet Union / Russia	*communism* is tied to state power: the United States faces "international communism," the Soviet Union becomes a communist "showpiece," and Russia appears as the site of Soviet economic law.
NORP	American / Soviet	National-political identities frame evaluation: American critics treat *communism* as a "nightmare," while Soviet writers claim that only communism can realize "true freedom."
ORG/GPE	Communist Party / Britain	Party institutions anchor the concept in debates over Marxism-Leninism, democratic centralism, and the name communism.
GPE	China / East Germany / Moscow	Geopolitical references place *communism* inside international competition, Soviet support, and national variants of communist rule.

Table 5: Representative concept–entity co-occurrences. NER shows that *communism* is anchored in national spaces, political identities, and party institutions rather than appearing as an abstract term alone.

5 BERT Contextual Frames

I embedded every sentence containing *communism* using local `bert-base-uncased`, with mean pooling over final hidden states. I then clustered the sentence embeddings with KMeans and visualized them with PCA. The BERT analysis was fully unsupervised: no weak labels were used to train, color, or name the clusters.

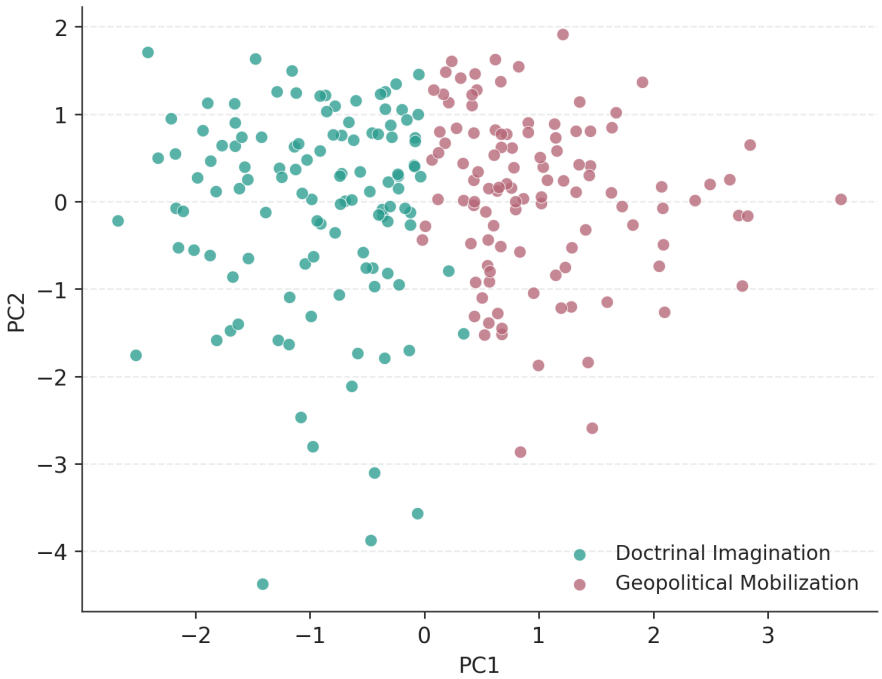

Figure 3: PCA projection of BERT sentence embeddings, colored by unsupervised cluster. The two groups were named only after reading representative examples, not from weak labels.

After close reading the representative examples, I named the two clusters DOCTRINAL IMAGINATION and GEOPOLITICAL MOBILIZATION. The split is not simply positive versus negative. Instead, it is an ontological split: it separates what kind of object *communism* becomes in context.

BERT Contextual Frame	Sentences	Share
DOCTRINAL IMAGINATION	117	51.5%
GEOPOLITICAL MOBILIZATION	110	48.5%

Table 6: Distribution of sentences across the two interpreted BERT contextual frames. The near balance is substantively important: the corpus does not overwhelmingly treat *communism* as only a doctrine or only a geopolitical force. Instead, it repeatedly alternates between imagining *communism* as a worldview and mobilizing it as a political actor.

This is the first major finding. BERT does not simply rediscover document buckets or stance labels. It finds a deeper distinction. In DOCTRINAL IMAGINATION, *communism* is a worldview to be explained, believed, doubted, or morally judged. In GEOPOLITICAL MOBILIZATION, it is a force that moves through states, parties, revolutions, and international strategy.

Frame	Text	Representative Excerpt
DOCTRINAL IMAGINATION	The Religion of Communism	communism is compared with historic religious movements and faith in a higher power
DOCTRINAL IMAGINATION	Nixon study paper	discussion of the "fundamental belief" of Communist economic philosophy
DOCTRINAL IMAGINATION	The Religion of Communism	communism has "metaphysical pretensions" and faith in the triumph of the proletarian cause
DOCTRINAL IMAGINATION	How to Tell a Communist	the fallacy of Communism lies not in goals but in practical plans for realizing them
GEOPOLITICAL MOBILIZATION	Eurocommunism is Anti-Communism	discussion of mass parties, capitalism, and changes in Western communism
GEOPOLITICAL MOBILIZATION	Voice of America anniversary speech	countries must be free of the "menace of international Communism"
GEOPOLITICAL MOBILIZATION	If America Should Go Communist	American communism is imagined as spreading to other countries
GEOPOLITICAL MOBILIZATION	Eurocommunism is Anti-Communism	Chinese leadership, American strategy, and postwar capitalist hegemony

Table 7: Representative examples from the BERT contextual frames. DOCTRINAL IMAGINATION treats *communism* as doctrine, belief, philosophy, or moral worldview. GEOPOLITICAL MOBILIZATION places it in parties, states, nations, movements, and international conflict.

6 Weak Supervision: Stance as a Second Layer

The BERT frames describe what kind of object *communism* becomes. But they do not directly tell us how a sentence evaluates that object. I therefore used weak supervision to create three stance labels: HOSTILE, NEUTRAL, and SYMPATHETIC. All three labels exceed the required minimum of 25 examples.

Weak Stance Label	Count
HOSTILE	102
NEUTRAL	79
SYMPATHETIC	46

Table 8: Distribution of weak stance labels. These labels capture how each sentence positions *communism* evaluatively.

The stance labels immediately reveal why sentence-level analysis matters. Corpus buckets are strongly associated with stance, but not identical to it. Official anti-communist texts are mostly hostile, mainstream commentary is mostly neutral, and left-internal critique is mostly sympathetic. Yet each bucket contains exceptions, showing that stance is a local property of sentence context rather than simply a document label.

Bucket	HOSTILE	NEUTRAL	SYMPATHETIC
Mainstream commentary	24	64	6
Official anti-communist	68	15	8
Western left-internal critique	10	0	32

Table 9: Stance label distribution by corpus bucket. The table confirms expected corpus-level tendencies while also showing that bucket membership does not fully determine sentence-level stance.

The more interesting pattern appears when stance is crossed with BERT frames. DOCTRINAL IMAGINATION is dominated by neutral and hostile sentences. GEOPOLITICAL MOBILIZATION is more polarized: it contains many hostile sentences, but also most of the sympathetic sentences. This means that political-geopolitical language is not automatically hostile. It is also the language of left-internal mobilization.

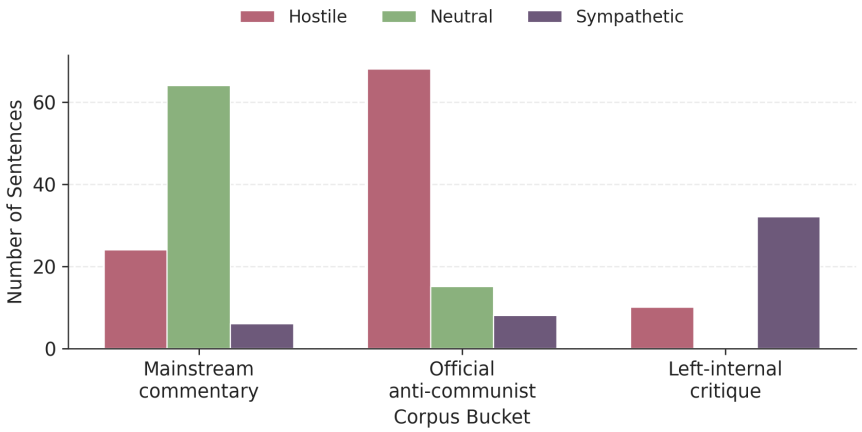

Figure 4: Weak stance labels by corpus bucket. Official anti-communist discourse concentrates hostile uses; mainstream commentary concentrates neutral explanation; left-internal critique concentrates sympathetic or insider framings.

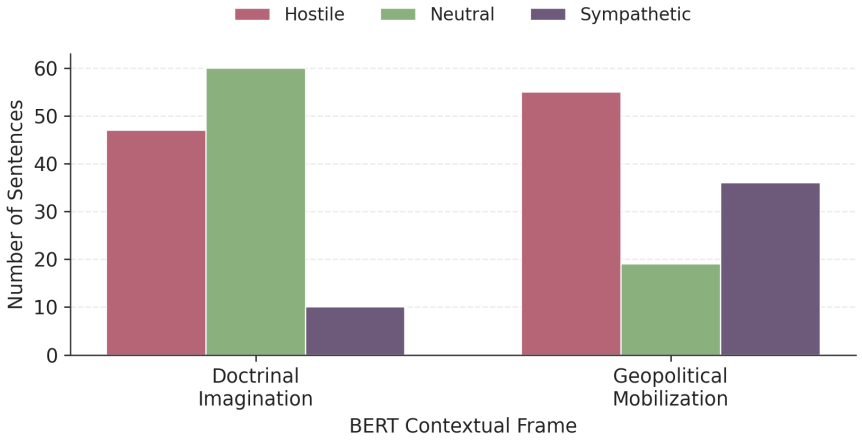

Figure 5: Weak stance labels across BERT contextual frames. DOCTRINAL IMAGINATION is more explanatory and moral-philosophical; GEOPOLITICAL MOBILIZATION is more polarized, containing both hostile Cold War framing and sympathetic left-internal mobilization.

BERT Frame	HOSTILE	NEUTRAL	SYMPATHETIC
DOCTRINAL IMAGINATION	47	60	10
GEOPOLITICAL MOBILIZATION	55	19	36

Table 10: Weak stance labels across BERT contextual frames. Contextual frame and stance are related, but not identical.

7 Logistic Regression Classification

To test whether stance can be predicted from local lexical evidence, I trained a TF–IDF logistic regression classifier with a 75/25 stratified train/test split. The model reaches 0.614 accuracy and 0.556 macro-F1. This is not a perfect classifier, but the pattern of performance is informative. The model distinguishes HOSTILE and NEUTRAL reasonably well, but struggles with SYMPATHETIC.

Label	Prec.	Rec.	F1	Support
HOSTILE	0.643	0.692	0.667	26
NEUTRAL	0.636	0.700	0.667	20
SYMPATHETIC	0.429	0.273	0.333	11
Accuracy	–	–	0.614	57
Macro avg.	0.569	0.555	0.556	57
Weighted avg.	0.599	0.614	0.602	57

Table 11: Logistic regression classification report for the three stance labels. The weaker performance on SYMPATHETIC reflects class imbalance and overlap between left-internal critique and hostile political vocabulary.

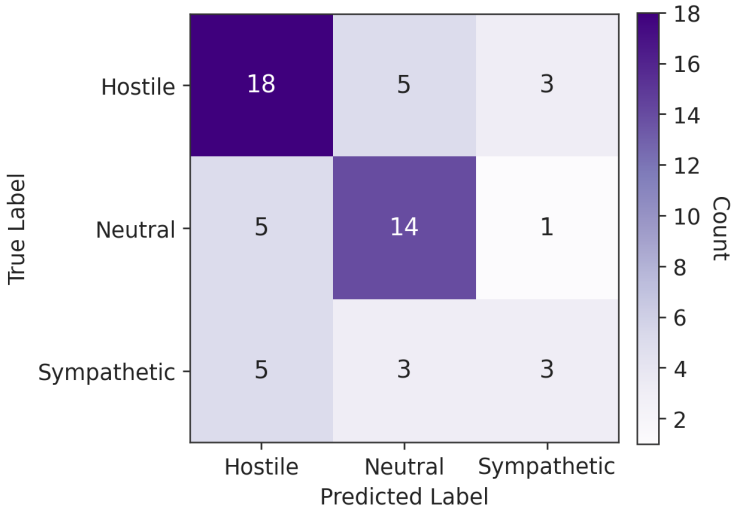

Figure 6: Confusion matrix for TF–IDF logistic regression. The model performs best on HOSTILE and NEUTRAL examples, while SYMPATHETIC is frequently confused with the other labels.

The top predictive features make the classifier interpretable. HOSTILE is driven by terms such as *international communism*, *war*, *freedom*, *tyranny*, *fight*, *false*, and *destroy*. NEUTRAL is driven by explanatory or conceptual terms such as *religion*, *philosophy*, *faith*, *scientific*, *theory*, and *history*. SYMPATHETIC is driven by vocabulary of political organization and revolutionary transformation, including *democracy*, *party*, *proletariat*, *revolution*, *capitalism*, and *socialism*.

The errors are not just failures; they clarify what makes the task difficult. A sentence such as "communism is utterly wrong about its most basic premise" is hostile, but the model predicts neutral because the vocabulary is philosophical. A left-internal sentence about the Soviet Union may be sympathetic in context but predicted hostile because it contains geopolitical terms associated with anti-communist discourse. This shows that stance often depends on document context and rhetorical position, not only local words.

Label	Top Positive Features
HOSTILE	international communism; war; freedom; law; international; challenge; tyranny; fight; false; destroy
NEUTRAL	religion; people; philosophy; support; faith; world communism; scientific; western; science; theory
SYMPATHETIC	democracy; party; proletariat; countries; revolution; Britain; population; capitalism; work; socialism

Table 12: Top positive logistic regression features by stance label. These features reveal how stance becomes lexical: hostile framing is security-oriented, neutral framing is explanatory, and sympathetic framing is organized around revolutionary and left-political vocabulary.

8 Synthesis: Two Dimensions of Meaning

The analyses converge on a two-dimensional account of *communism*. The first dimension is contextual ontology: what kind of object is *communism* imagined to be? BERT answers this with two frames. In DOCTRINAL IMAGINATION, *communism* is a doctrine, worldview, faith, philosophy, or moral challenge. In GEOPOLITICAL MOBILIZATION, it is a movement, party, state project, geopolitical force, or international actor.

The second dimension is stance: how does the sentence position that object? Weak supervision and logistic regression show three stance positions: HOSTILE, NEUTRAL, and SYMPATHETIC. These stances are related to corpus bucket, but not reducible to it. They also interact with BERT frames. DOCTRINAL IMAGINATION tends toward neutral explanation and moral judgment, while GEOPOLITICAL MOBILIZATION becomes the site of both Cold War hostility and left-internal mobilization.

The central claim is therefore not that Western discourse has one dominant meaning of *communism*. Rather, Western discourse constructs *communism* through the interaction of imagination and mobilization. It first decides what kind of thing communism is: a doctrine, a philosophy, a religion-like worldview, a party, a state project, or an international force. It then evaluates that object as dangerous, explainable, or politically desirable.

9 Limitations

Several limitations remain. First, the unit of analysis is the sentence, which makes modeling tractable but can miss document-level irony, quotation, and argument structure. Second, the extraction focuses on the exact word *communism*; related forms such as *communist*, *Marxism*, and *socialism* are not included unless they occur in the same sentence. Third, the stance labels are weak labels, so they are useful but noisy. This is especially true for left-internal critique, where negative language can be used in defense of communism as an ideal. Fourth, spaCy NER is imperfect for historical-political prose and can miss abstract political objects such as *free world*, *market economy*, or *class struggle*. Finally, BERT clustering should be interpreted as a useful structure, not as a natural taxonomy.

10 Conclusion

This project shows that *communism* is not a stable keyword with one meaning. Across the corpus, it is constructed through two broad contextual frames and three stance positions. BERT reveals a distinction between DOCTRINAL IMAGINATION and GEOPOLITICAL MOBILIZATION. Weak supervision then shows that each frame can be evaluated through hostile, neutral, or sympathetic stance. Together, these methods

show how a contested political concept becomes meaningful: not by carrying one fixed definition, but by moving between worldview and mobilization, explanation and judgment, doctrine and geopolitical force.

References

Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL-HLT*, 2019.

Maud Ehrmann, Matteo Romanello, Alex Flückiger, and Simon Clematide. Named entity recognition and classification in historical documents: A survey. *ACM Computing Surveys*, 2023.

J. R. Firth. *A Synopsis of Linguistic Theory, 1930–1955*. Blackwell, 1957.

Justin Grimmer and Brandon M. Stewart. Text as data: The promise and pitfalls of automatic content analysis methods for political texts. *Political Analysis*, 21(3):267–297, 2013.

William L. Hamilton, Jure Leskovec, and Dan Jurafsky. Diachronic word embeddings reveal statistical laws of semantic change. In *Proceedings of ACL*, pages 1489–1501, 2016.

Austin C. Kozlowski, Matt Taddy, and James A. Evans. The geometry of culture: Analyzing the meanings of class through word embeddings. *American Sociological Review*, 84(5):905–949, 2019.

Guillaume Lample, Miguel Ballesteros, Sandeep Subramanian, Kazuya Kawakami, and Chris Dyer. Neural architectures for named entity recognition. In *Proceedings of NAACL-HLT*, pages 260–270, 2016.

Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*, 2013.

Jeffrey Pennington, Richard Socher, and Christopher D. Manning. GloVe: Global vectors for word representation. In *Proceedings of EMNLP*, pages 1532–1543, 2014.

Alexander Ratner, Stephen H. Bach, Henry Ehrenberg, Jason Fries, Sen Wu, and Christopher Ré. Snorkel: Rapid training data creation with weak supervision. *The VLDB Journal*, 29:709–730, 2020.

Margaret E. Roberts, Brandon M. Stewart, Dustin Tingley, Christopher Lucas, Jetson Leder-Luis, Shana Kushner Gadarian, Bethany Albertson, and David G. Rand. Structural topic models for open-ended survey responses. *American Journal of Political Science*, 58(4):1064–1082, 2014.